

COMP_SCI 214: Data Structures and Algorithms

Single-Source Shortest Path

PROF. SRUTI BHAGAVATULA

Announcements

- ▶ Homework 5 due today
- ▶ Homework 4 self-eval due today
- ▶ Project to be released later today
 - ▶ Much larger than a homework
 - ▶ and hence much more time and more submissions
 - ▶ START EARLY!!!
 - ▶ Details forthcoming on Piazza and in the handout --- Read fully!


Another class of graph algorithms

Problem of the day

- ▶ Given a graph containing locations (nodes) connected via roads (edges) of different lengths (weights), how can we find the shortest path between two nodes?
- ▶ Basically: how does Google Maps do this?

The problem: Finding shortest paths


- On this weighted graph, what is the shortest path from A to E?



The problem: Finding shortest paths

- On this weighted graph, what is the shortest path from A to E?


- A-D-B-E: 14



The problem: Finding shortest paths

- ▶ On this weighted graph, what is the shortest path from A to E?


- ▶ A-D-B-E: 14
- ▶ A-B-E: 15. Worse!



The problem: Finding shortest paths

- ▶ On this weighted graph, what is the shortest path from A to E?

- ▶ A-D-B-E: 14
- ▶ A-B-E: 15. Worse!
- ▶ A-D-B-C-E: 10
 - ▶ We have a winner!



The problem, generalized

- ▶ Find the shortest path to everywhere from a given vertex
 - ▶ **Single-source shortest path (SSSP)** problem
- ▶ Why generalize?
 - ▶ To find SP to one destination, we may end up finding the shortest path to other destinations anyway
 - ▶ So might as well
- ▶ Applies to directed and undirected graphs alike

Why not BFS?

- ▶ BFS can do exactly this
 - ▶ If #edges in a path is the path length
 - ▶ Because you are always minimizing number of edges in path by checking paths of a specific length before moving onto next level of neighbors
- ▶ But what about when we have weights?
 - ▶ Path length = sum of weights in path
 - ▶ BFS doesn't consider this
 - ▶ Need to modify BFS or use a different approach

Problem setup

- ▶ For each vertex in the graph, we want to find:
 - ▶ The minimum “cost” of getting from starting vertex to vertex in question
 - ▶ The actual path that gives us this minimal cost
- ▶ Think of “cost” as the sum of weights along a path
- ▶ Let’s maintain the cost of reaching every vertex and the predecessor of the vertex in this minimal path

The solution: Just relax

- ▶ Relaxation:
 - ▶ Set cost for reaching each vertex from starting vertex to its maximum: ∞
 - ▶ Bring cost of reaching a vertex down as we learn more information about the rest of the graph
 - ▶ Once you've relaxed enough, you have the real cost of reaching a node
- ▶ General approach to algorithms


SSSP setup

- ▶ Let's maintain two pieces of information:
 1. The cost of reaching every vertex from the starting vertex ($\text{dist}[v]$)
 - ▶ We'll keep updating this as we learn about the graph
 2. The predecessor of the vertex in this minimal path ($\text{pred}[v]$)
 - ▶ At the end, we can trace this back to get the shortest path

Relaxation for SSSP

- ▶ SSSP from A
- ▶ To start, most pessimistic costs possible: everything ∞
- ▶ As we look at edges, we learn more information about the cost to get to vertices


v	dist[v]	pred[v]
A	∞	
B	∞	
C	∞	
D	∞	
E	∞	



Relaxation for SSSP

- ▶ Visit A first
- ▶ Since node is itself, path length is 0


v	dist[v]	pred[v]
A	0	
B	∞	
C	∞	
D	∞	
E	∞	



Relaxation for SSSP

- ▶ Next visit B via edge A-B (A's neighbor)
- ▶ Distance from A-B is 5
 - ▶ Update $\text{dist}[B]$ and $\text{pred}[B]$ in table


v	dist[v]	pred[v]
A	0	
B	5	A
C	∞	
D	∞	
E	∞	



Relaxation for SSSP

- ▶ Let's visit C via the B-C edge now (B's neighbor)
- ▶ Distance from A-C = A-B distance + B-C distance = $5 + 4 = 9$
 - ▶ Update $\text{dist}[C]$ and $\text{pred}[C]$ in table


v	dist[v]	pred[v]
A	0	
B	5	A
C	9	B
D	∞	
E	∞	



Relaxation for SSSP

- ▶ Let's visit E via the C-E edge now (C's neighbor)
- ▶ Distance from A-E = A-C distance + C-E distance = $9 + 2 = 11$
 - ▶ Update $\text{dist}[E]$ and $\text{pred}[E]$ in table


v	dist[v]	pred[v]
A	0	
B	5	A
C	9	B
D	∞	
E	<u>11</u>	C



Relaxation for SSSP

- ▶ Let's visit B via the E-B edge now (E's neighbor)
- ▶ Distance from A-B = A-E distance + E-B distance = $11 + 10 = 21$
 - ▶ Not as good as $\text{dist}[B] = 5$; don't update


v	dist[v]	pred[v]
A	0	
B	5	A
C	9	B
D	∞	
E	11	C



Relaxation for SSSP

- ▶ Let's now look at D via the A-D edge (A's neighbor from earlier)
- ▶ Distance from A-D = A-A distance + A-D distance = $0 + 1 = 1$
- ▶ Update `dist[D]` and `pred[D]`


v	dist[v]	pred[v]
A	0	
B	5	A
C	9	B
D	1	A
E	11	C



Relaxation for SSSP

- ▶ Let's now look at B via the D-B edge (D's neighbor)
- ▶ Distance from A-B = A-D distance + D-B distance = $1 + 3 = 4$
 - ▶ Better than old `dist[B]` so update in table


v	dist[v]	pred[v]
A	0	
B	5 4	A D
C	9	B
D	1	A
E	11	C



Relaxation for SSSP

- ▶ Oh, but our graph is undirected! $D-B = B-D$!
- ▶ Did we find a shorter path to D via B too? No. $5+3 > 1$ in the graph
 - ▶ Each edge direction needs to be treated separately even if undirected

v	dist[v]	pred[v]
A	0	
B	4	D
C	9	B
D	1	A
E	11	C




Relaxing an edge

- ▶ Given the starting vertex s
- ▶ Relaxing an edge (u, v) means checking if a path from s to v which passes through (u, v) is shorter than other paths to v seen before
 - ▶ To start, since all paths to vertices start at ∞ , any path will be shorter the first time we relax an edge
- ▶ If (u, v) makes the path shorter, we reduce the cost to reach v
 $\text{if } \text{dist}[u] + w(u, v) < \text{dist}[v]:$
 $\text{dist}[v] \leftarrow \text{dist}[u] + w(u, v)$

Relaxation for SSSP

- ▶ Are we done? Maybe! But we'll stop here.
- ▶ Different algorithms have different stopping conditions
 - ▶ All guaranteed to find shortest path to each node

v	dist[v]	pred[v]
A	0	
B	4	D
C	9	B
D	1	A
E	11	C



What do we need from an algorithm?

- ▶ Systematic way to relax edges further after a round of relaxing
- ▶ A clear stopping condition


Dynamic programming

- ▶ To determine shortest path between a and b , we use an intermediate shortest path on the way (e.g., a and q) and build on that
- ▶ These algorithms that solve sub-problems and determine the final solutions using the results of sub-solutions use **dynamic programming**
 - ▶ You should see more of this in later classes!

Bellman-Ford Algorithm

A note on negative path lengths

- ▶ SSSP doesn't make sense for all graphs.



- ▶ A-B-C-E has length 3.
- ▶ A-B-C-D-E has length -6. Shorter!
- ▶ A-B-C-D-E-C-D-E has length -13. Even shorter!


Negative cycles

- ▶ Negative cycle = path length with negative value
- ▶ Negative cycles means no shortest *est* path

Bellman-Ford algorithm

- ▶ **Solves:** SSSP for graphs with no negative cycles
- ▶ **Main idea:** Relax every edge $v - 1$ times
- ▶ **Time complexity:** ?
- ▶ **Steps:**
 - ▶ Repeat steps 1-2 $v - 1$ times
 1. Pick edges in some order
 2. Relax the edges

A	0
B	10
C	2
D	15 5



Initialization for Bellman-Ford algorithm

Input: A graph `graph` and starting vertex `start`

Output: Table of vertex distances `dist` and predecessors `pred`

```
for every vertex v in graph do
    dist[v] ← ∞; pred[v] ← None
end
```

```
dist[start] ← 0
```

v	dist[v]	pred[v]
0	∞	
1	∞	
2	∞	
3	∞	
...		

The Bellman-Ford algorithm

```
for |Vertices(graph)|-1 iterations do
  for every edge (v, u) with weight w in graph do
    # This conditional relaxes edge v-u.
    if dist[v] + w < dist[u] then
      # (if undirected, we'd also relax edge u-v)
      dist[u] ← dist[v] + w;
      pred[u] ← v;
    end
  end
end
end
```

Why does $v - 1$ iterations work?

- ▶ Iteration 1: Nodes whose shortest path is 1 edge away now have their correct shortest paths
 - ▶ No guarantees about the other nodes because of arbitrary order
- ▶ Iteration 2: Nodes whose shortest path is 2 edges away now have their correct shortest paths
- ▶ Iteration 3: Nodes whose shortest path is 3 edges away now have their correct shortest paths
- ▶ ...
- ▶ A shortest path between nodes can have a max of $v - 1$ edges, therefore we'll have the largest shortest path after $v - 1$ iterations

Bellman-Ford with negative cycles

- ▶ If there are no negative cycles, we're done with Bellman-Ford
- ▶ But we should check for negative cycles:
 - ▶ After we're done with Bellman-Ford, try to relax all the edges one more time
 - ▶ If a cost reduces further, we have a negative cycle
 - ▶ Shouldn't be possible to improve further after $v - 1$ iterations otherwise
 - ▶ There will never be a minimum then as cost can keep reducing
- ▶ Are we satisfied to go to use Bellman-Ford for our map problem now?

Revisit: problem of the day

- ▶ Given a graph containing locations (nodes) connected via roads (edges) of different lengths (weights), how can we find the shortest path between two nodes?
- ▶ This graph may have LOTS of edges in $O(v^2)$ in the worst case if it's a dense graph

The issue with Bellman-Ford

- ▶ Bellman-Ford is $O(e * v)$
- ▶ ...which, in the worst case, becomes $O(v^3)$ if edges in $O(v^2)$
 - ▶ Cubic, yikes!

Wastage in Bellman-Ford

- ▶ We may not need all $v - 1$ iterations of relaxations
 - ▶ Could stop when a round of relaxing all edges doesn't improve anything
- ▶ Relaxing edges in an arbitrary order may not be optimal
 - ▶ A more clever order could get us there faster
 - ▶ Let's see another algorithm that does this; relaxes each edge *once*!

Dijkstra's Algorithm

Dijkstra's algorithm


- ▶ Named after Edsger W. Dijkstra (pronounced: “dike-strah”)
 - ▶ Early CS pioneer, 1972 Turing award
 - ▶ A name you'll see often in CS

Dijkstra's algorithm

- ▶ **Solves:** SSSP for graphs with **<some restrictions>**
- ▶ **Main idea:** Relax the edges in a clever order (like BFS)
- ▶ **Time complexity:** ?
- ▶ Main approach:
 - ▶ Relax the edges coming out of the *nearest* vertex so far
 - ▶ Then repeat with next nearest, etc.
 - ▶ Keep track of edges already relaxed; relax each edge only once!

Dijkstra's algorithm in action: SSSP from A

v	dist	pred
A	0	
B	∞	
C	∞	
D	∞	
E	∞	
F	∞	




Left: {A, B, C, D, E, F}

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	∞ 7	
C	∞	
D	∞	
E	∞ 14	
F	∞ 9	

Step:

A has lowest `dist` in table so far, visit it




Left: {A, B, C, D, E, F}

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	∞	
D	∞	
E	14	A
F	9	A

Step:

For each of A's neighbors n , relax (A, n)




Left: ~~{A, B, C, D, E, F}~~

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	∞	
D	∞	
E	14	A
F	9	A

Step:

Next, look at A's nearest unvisited vertex: B



Left: {B, C, D, E, F}


Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	22	B
D	∞	
E	14	A
F	9	A

Now that we're at B, we **know** we found the shortest path to it

Step:

For each of B's neighbors n , relax (B, n)




Left: ~~{B, C, D, E, F}~~

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	22	B
D	∞	
E	14	A
F	9	A

Step:

Next, look at A's nearest unvisited vertex: F



Left: {C, D, E, F}


Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	∞	
E	11	F
F	9	A

Now that we're at F, we **know** we found the shortest path to it

Step:

For each of F's neighbors n , relax (F, n)




Left: {C, D, E, F}

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	∞	
E	11	F
F	9	A

Step:

Next, look at A's nearest unvisited vertex: E



Left: {C, D, E}


Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	20	E
E	11	F
F	9	A

Now that we're at E, we **know** we found the shortest path to it

Step:


For each of E's neighbors n , relax (E, n)



Left: {C, D, ~~E~~}

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	20	E
E	11	F
F	9	A

**Step:**

Next, look at A's nearest unvisited vertex: C

Left: {C, D}


Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	20	E
E	11	F
F	9	A

Now that we're at C, we **know** we found the shortest path to it

Step:

For each of C's neighbors n , relax (C, n)




Left: {C, D}

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	20	E
E	11	F
F	9	A

Step:

Next, look for A's nearest unvisited vertex: D



Left: {D}


Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	20	E
E	11	F
F	9	A

Now that we're at D, we **know** we found the shortest path to it

Step:


For each of D's neighbors n , relax (D, n)



Left: {D}

Dijkstra's algorithm in action

v	dist	pred
A	0	
B	7	A
C	20	F
D	20	E
E	11	F
F	9	A



No vertices left to visit, we found all the shortest paths!

Left: {}

Recovering the shortest path

- ▶ Start from destination and build path backwards using pred

- ▶ Just like in DFS/BFS!

- ▶ **Quiz:** Shortest path from A to D?

1. A-E-D

2. A-F-E-D

3. A-B-C-D

A - F - E - D

v	dist	pred
A	0	
B	7	A
C	20	F
D	20	E
E	11	F
F	9	A

Initialization for Dijkstra's algorithm v1

Input: A graph `graph` and starting vertex `start`

Output: Table of vertex distances `dist` and predecessors `pred`

```
for every vertex v in graph do
    dist[v] ← ∞; pred[v] ← None
end
```

```
dist[start] ← 0
left ← set of vertices in graph
```

v	dist	pred
0	∞	
1	∞	
2	∞	
3	∞	
...		

Dijkstra's algorithm v1

left \leftarrow set of vertices in graph

```
while left is not empty do
  v  $\leftarrow$  remove the element of left with minimal dist[v];
  for every outgoing edge (v,u) with weight w do
    if dist[v] + w < dist[u] then
      dist[u]  $\leftarrow$  dist[v] + w;
      pred[u]  $\leftarrow$  v;
    end
  end
end
```

Dijkstra's algorithm v1

`left` \leftarrow set of vertices in graph

```
while left is not empty do
  v  $\leftarrow$  remove the element of left with minimal dist[v];
  for every outgoing edge (v,u) with weight w do
    if dist[v] + w < dist[u] then
      dist[u]  $\leftarrow$  dist[v] + w;
      pred[u]  $\leftarrow$  v;
    end
  end
end
```

Cool algorithm.

But how do you find the minimal element in `left`?

We need a worklist

- ▶ BFS used a queue, but we don't want FIFO
- ▶ We want the minimum element in a worklist
 - ▶ Or element of highest priority?
- ▶ We can use a priority queue!
 - ▶ Values are vertices to look at next (like in BFS)
 - ▶ Priorities are $\text{dist}[v]$

Recall: Priority queue ADT

- ▶ Abstract values look like (note the sorting) → Highest priority
- ▶ Priority-value pairs

2	Brain damage
5	Heart attack
17	Fever
89	Cold
...	...

```
interface PRIORITY_QUEUE[T]:  
  def empty?(self) -> bool?  
  def insert(self, priority: num?, value: T) -> NoneC  
  def remove_min(self) -> T
```

Initialization for Dijkstra's algorithm v2 w/ Priority Queue

Input: A graph `graph` and starting vertex `start`

Output: Table of vertex distances `dist` and predecessors `pred`

```
for every vertex v in graph do
    dist[v] ← ∞; pred[v] ← None
end
```

```
dist[start] ← 0
```

```
todo ← empty priority queue;
```

```
done ← empty vertex set e.g., mark array
```


```
todo.insert(0, start);
```

Dijkstra's algorithm v2 w/ Priority Queue

```
while todo is not empty do
  v ← todo.remove_min(); #pick the nearest vertex
  if v ∉ done then
    done ← done ∪ {v};
    for every outgoing edge (v,u) with weight w do
      #relax outgoing edges
      if dist[v] + w < dist[u] then
        dist[u] ← dist[v] + w;
        pred[u] ← v;
        todo.insert(dist[u], u);
      end
    end
  end
end
```


In-class exercise (6 minutes)

1. Run Dijkstra's algorithm *from A* and write the results of the table in the format on the right
2. Dijkstra's algorithm cannot find shortest path in this graph from A-C. Why might that be based on your table above?




A, $d[v]$, $p[v]$
B, $d[v]$, $p[v]$
C, $d[v]$, $p[v]$

v	dist[v]	pred[v]
A	0	
B		
C		

In-class exercise solution

1. Actual shortest path: A-B-C with path length of $4 - 2 = 2$
2. Dijkstra's algorithm would not be able to find this path because it would prioritize first going to C from A and then would relax the C-B edge overwriting `dist[B]`
 - ▶ Dijkstra's wouldn't be able to find a path from A-C as following the above approach it would give us C-B-C-B-C... by following predecessors from C
 - ▶ No guarantees on an optimal path or a path at all
- ▶ Hence...



v	dist[v]	pred[v]
A	0	
B	1	C
C	-1	B

Dijkstra's algorithm

- ▶ **Solves:** SSSP for graphs with **no negative edge weights**
- ▶ **Main idea:** Relax the edges in a clever order (like BFS)
- ▶ **Time complexity:** ?
- ▶ More restrictive but a more optimal algorithm

Time complexity of Dijkstra's algorithm

```
while todo is not empty do
  v ← todo.remove_min(); #pick the nearest vertex
  if v ∉ done then
    done ← done ∪ {v};
    for every outgoing edge (v,u) with weight w do
      #relax outgoing edges
      if dist[v] + w < dist[u] then
        dist[u] ← dist[v] + w;
        pred[u] ← v;

        todo.insert(dist[u], u);
      end
    end
  end
end
```

Time complexity of Dijkstra's algorithm

- ▶ Relax every edge once: $O(e)$
- ▶ For every edge we relax, we do an `insert`
 - ▶ Which takes $O(\log e)$
- ▶ For every edge we relax, we do a `remove_min`
 - ▶ Which takes $O(\log e)$
- ▶ So Dijkstra's algorithm is $O(e \log e)$
 - ▶ Which is bounded by $O(v^2 \log v^2) = O(v^2 \log v)$ in the worst case
 - ▶ Better than Bellman-Ford's $O(v^3)$